

Virtual Private Network

Introduction

- Networks primarily intended for internal use are called private network.
- If we grant access from outside to the private network, the attack surface will significantly broaden.
- If the internal resources still use IP address as the basis for authorization, it is not difficult for attackers to access the protected resources


Virtual Private Network

VPN allows users to create a secure, private network over a public network such as the Internet. This is achieved by:

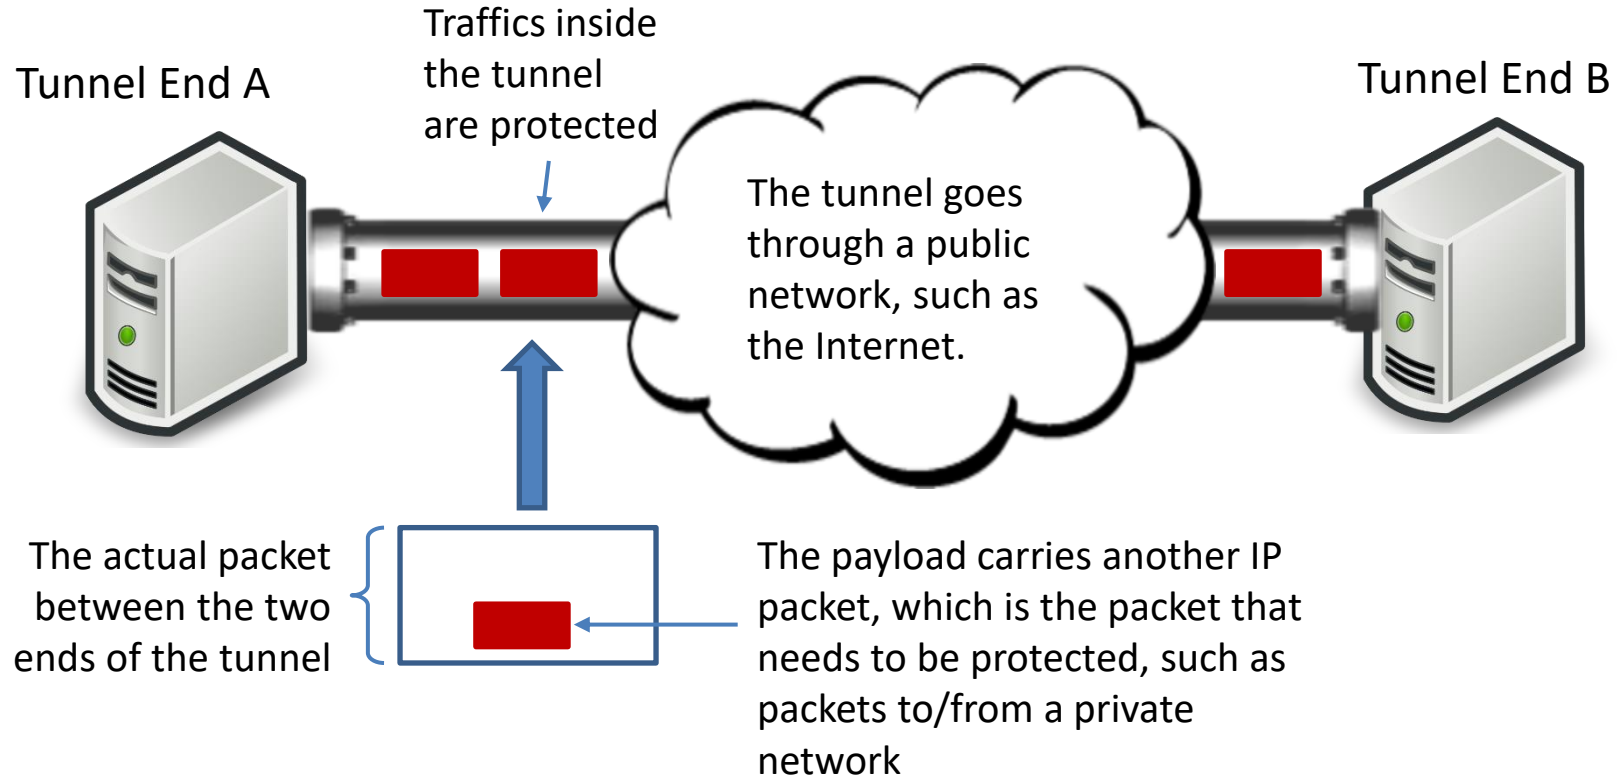
- Having a designated host (VPN server) on the network
- Outside computers have to go through the VPN server to reach the hosts inside a private network via authentication.
- VPN server is exposed to the outside and the internal computers are still protected, via firewalls or reserved IP addresses.

A Typical Setup

This is a typical VPN setup where the “Client” machine wants to connect with machine “V” on a private network. “Client” uses the “VPN Server” to get authenticated to the private network



IP Tunneling

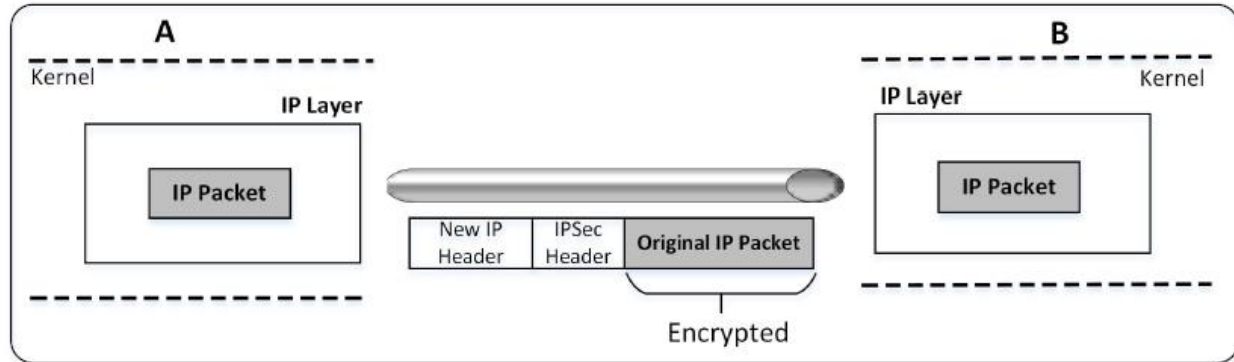


Two Types of IP Tunneling

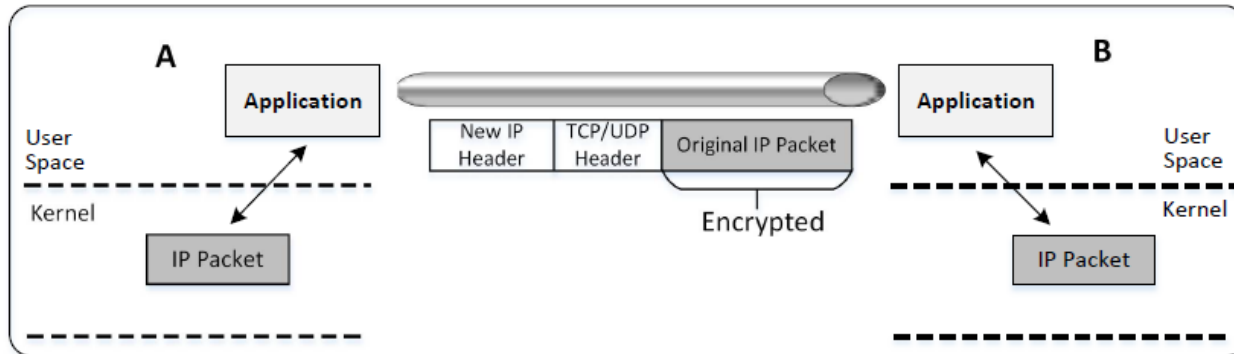
- IPSec Tunneling:
 - Utilizes the Internet Protocol Security protocol
 - IPSec has a mode called Tunneling mode, where the original IP packet is encapsulated and placed into a new IP packet
- TLS/SSL Tunneling:
 - Tunneling done outside the kernel, at the application level
 - Idea is to put each VPN-bound IP packet inside a TCP or UDP packet
 - The other end of the tunnel will extract the IP packet from the TCP/UDP payload
 - To secure the packets, both ends will use TLS/SSL protocol on top of TCP/UDP

Two Types of IP Tunneling

IPSec Tunneling



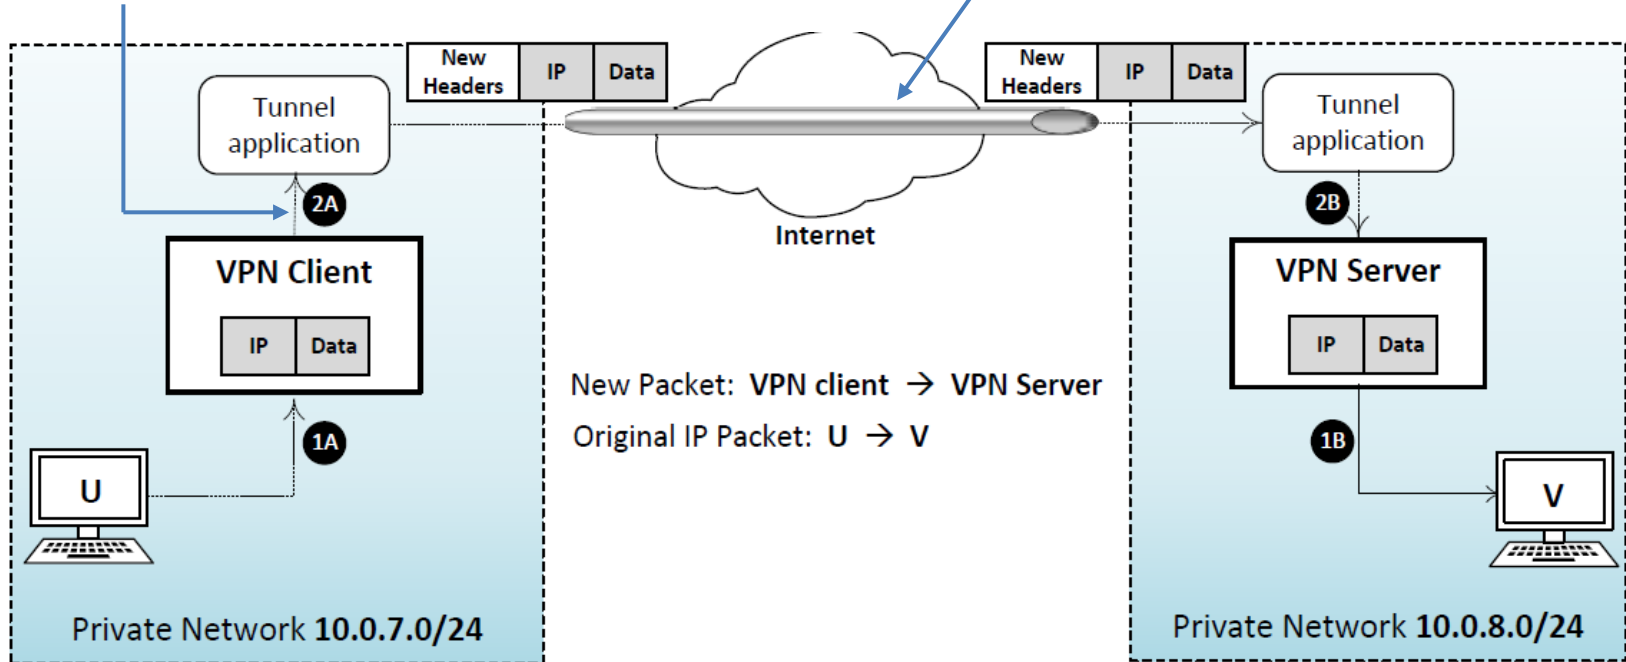
TLS/SSL Tunneling
(we will focus on
this type)



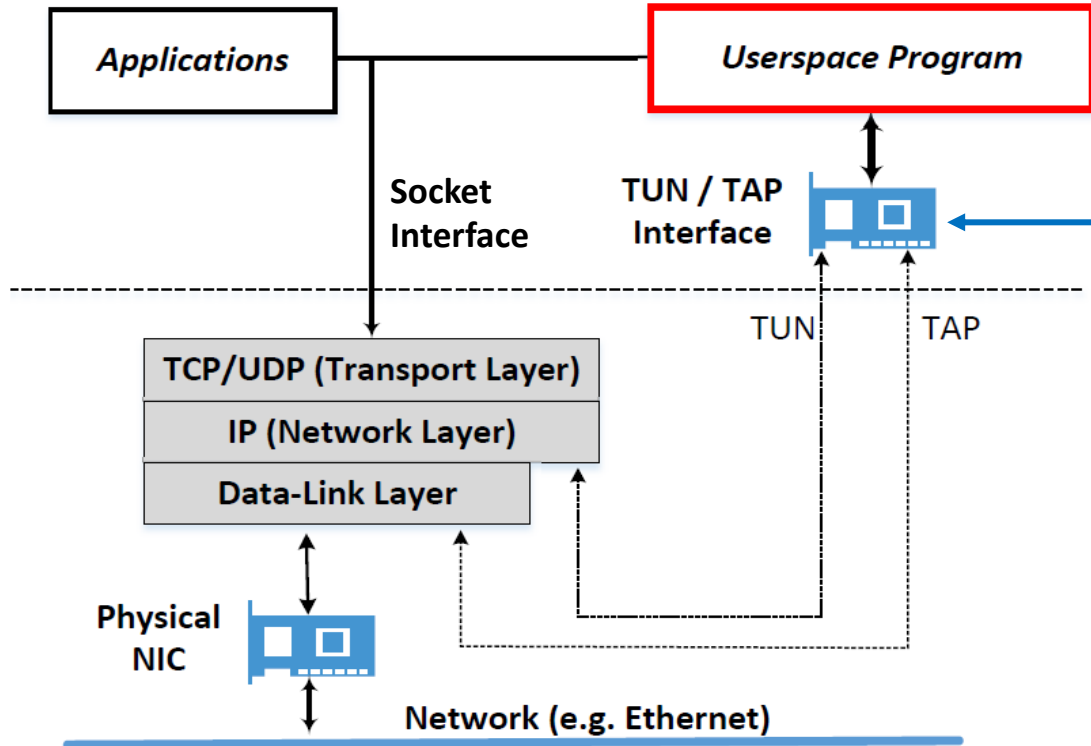
An Overview of How TLS/SSL VPN Works

Question: How can the Tunnel application get an IP packet?

This is just a normal TCP or UDP based SSL connection



TUN/TAP Interface



- **Question:** How can the Tunnel application get an IP packet?
 - Typically, applications interact with kernel using socket
 - Using socket, kernel only gives the data part of a packet to applications
 - Applications need to use a different way to interact with kernel

TUN/TAP Interface

- Most operating systems have two types of network interfaces:
 - Physical: Corresponds to the physical Network Interface Card (NIC)
 - Virtual: It is a virtualized representation of computer network interfaces that may or may not correspond directly to the NIC card. Example: *loopback* device
- TUN Virtual Interface
 - Work at OSI layer 3 or IP level
 - Sending any packet to TUN will result in the packet being delivered to user space program
- TAP Virtual Interfaces
 - Work at OSI layer 2 or Ethernet level
 - Used for providing virtual network adapters for multiple guest machines connecting to a physical device of the host machine

Creating a TUN Interface

```
int tunfd;  
struct ifreq ifr;  
memset(&ifr, 0, sizeof(ifr));  
  
ifr.ifr_flags = IFF_TUN | IFF_NO_PI; ①  
  
tunfd = open("/dev/net/tun", O_RDWR); ②  
ioctl(tunfd, TUNSETIFF, &ifr); ③
```

The flag IFF_TUN specifies that we are creating a TUN interface



Configure the TUN Interface

- Find the TUN interface

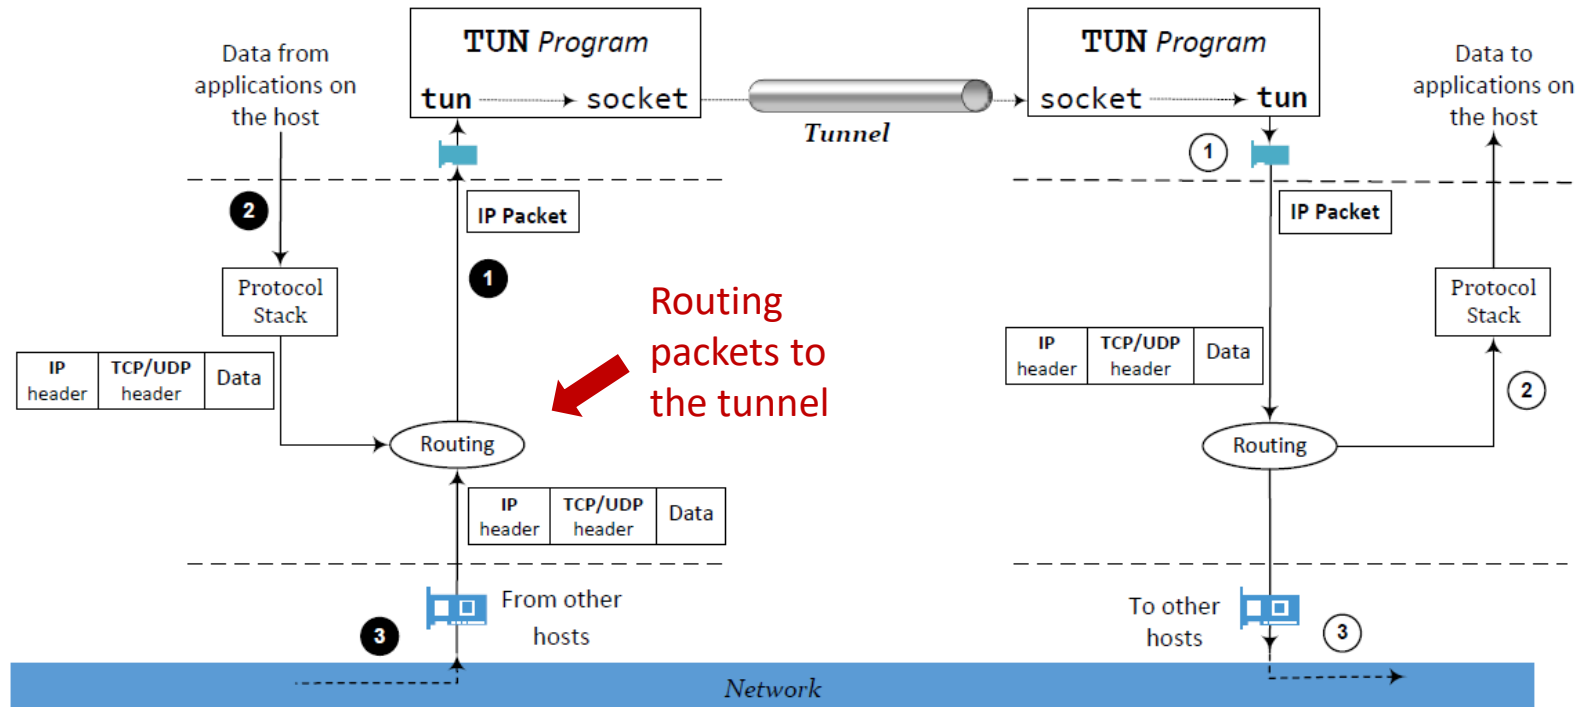
```
% ifconfig -a
tun0  Link encap:UNSPEC  HWaddr 00-00-00 ...
      POINTOPOINT NOARP MULTICAST  MTU:1500 ...
```

- Assign an IP address to the TUN interface and bring it up

```
% sudo ifconfig tun0 10.0.8.99/24 up

% ifconfig
tun0 Link encap:UNSPEC  HWaddr 00-00-00 ...
     inet addr: 10.0.8.99 P-t-P:10.0.8.99 Mask: 255.255.255.0
     UP POINTOPOINT RUNNING NOARP MULTICAST  MTU:1500 ...
```

Set UP the Routing

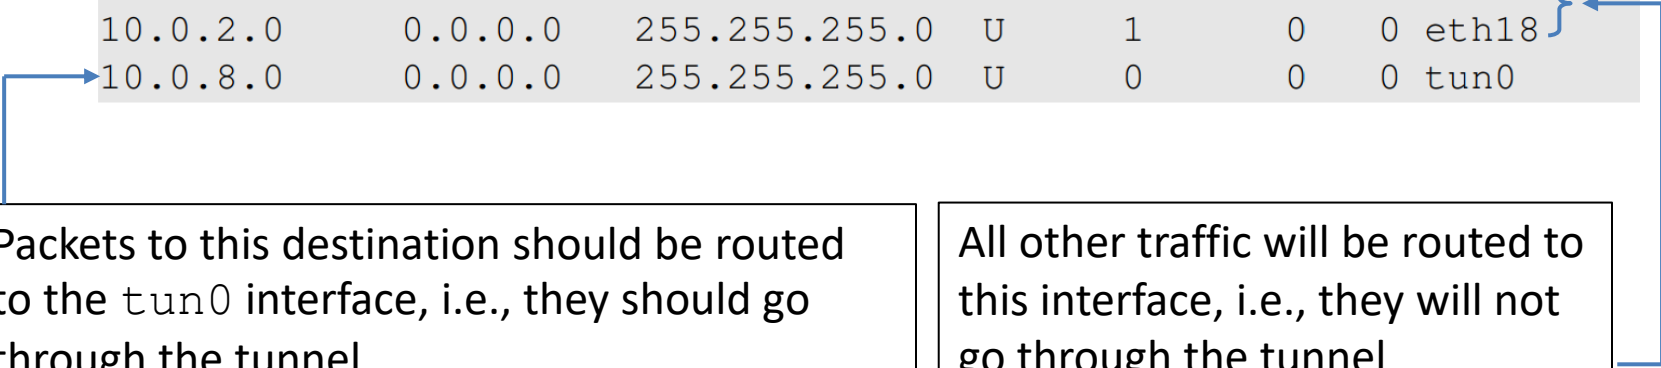


Set UP the Routing

```
$ sudo route add -net 10.0.8.0/24 tun0
```

```
$ route -n
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	10.0.2.1	0.0.0.0	UG	0	0	0	eth18
10.0.2.0	0.0.0.0	255.255.255.0	U	1	0	0	eth18
10.0.8.0	0.0.0.0	255.255.255.0	U	0	0	0	tun0



Packets to this destination should be routed to the `tun0` interface, i.e., they should go through the tunnel.

All other traffic will be routed to this interface, i.e., they will not go through the tunnel

Experiment: Reading From TUN Interface

We did an experiment by sending a ping packet to 10.0.8.32. The packet was sent to the TUN interface and then to our program. We use “xxd” to read from the interface and convert the into hexdump.

IP Header

```
$ sudo ./tundemo
```

```
TUN file descriptor: 3
```

0a00 0863: Source IP (128.0.8.99)

```
# xxd <& 3
```

```
00000000: 4500 0054 0000 4000 4001 1627 0a00 0863  E..T..@.@...'...c
00000010: 0a00 0820 0800 3b19 10cf 0001 da1d 9f57  ... ..;.....W
00000020: 439e 0400 0809 0a0b 0c0d 0e0f 1011 1213  C.....
00000030: 1415 1617 1819 1a1b 1c1d 1e1f 2021 2223  ..... !"#
00000040: 2425 2627 2829 2a2b 2c2d 2e2f 3031 3233  $%&'()*+,-./0123
00000050: 3435 3637
```

0a00 0820: Destination IP (128.0.8.32)

Experiment: Writing To TUN Interface

- We can write data to TUN interfaces.
- We can create a valid packet using the same “xxd” command.
- Copy-paste the xxd output from the previous slide into a file called “hexfile” and run “xxd -r hexfile > packetfile”.
- Now we write the packetfile to the interface:


```
# cat packetfile >& 3
```

- We should be able to observe the packet using Wireshark.

Establish a Transport-Layer Tunnel

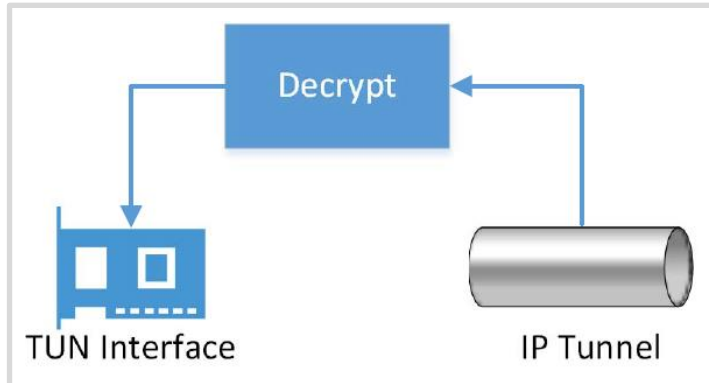
- A tunnel is just a TLS/SSL connection.
- Two applications (VPN client and server applications) just establish a TLS/SSL connection between themselves.
- Traffics inside are protected by TLS/SSL
- What makes this TLS/SSL connection a tunnel?
 - The payloads inside are IP packets
 - That is why it is called IP tunnel

How to Send/Receive Packets via Tunnel



Sending a packet via the tunnel

- Get an IP packet from the TUN interface
- Encrypt it (also add MAC)
- Send it as a payload to the other end of the tunnel

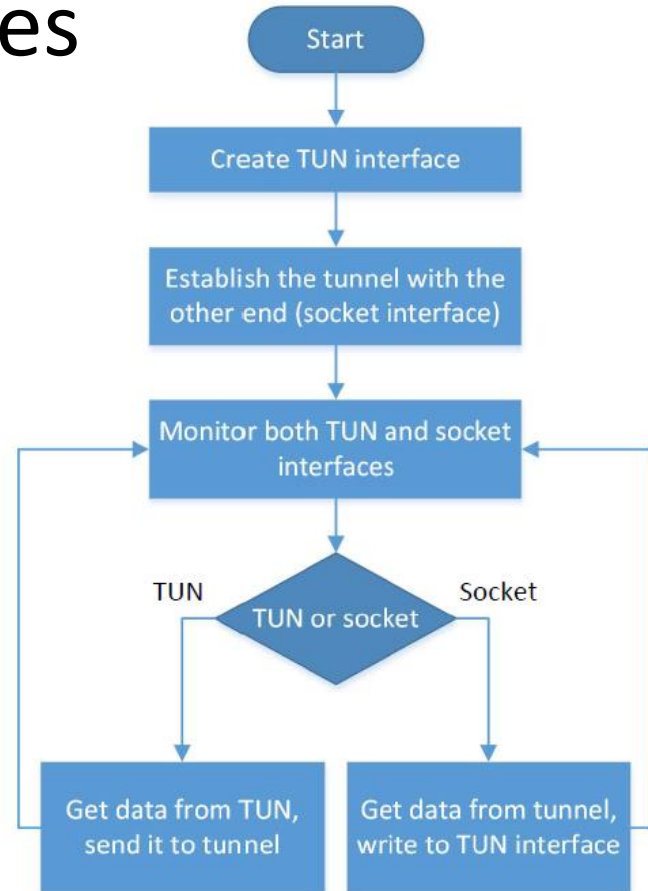


Receiving a packet from the tunnel

- Get a payload from the tunnel
- Decrypt it and verify its integrity
- We get the actual packet
- Write the packet to the TUN interface


Monitoring Both Interfaces

- Each tunnel application has two interfaces: socket and TUN
- Need to monitor both
- Forward packets between these two interfaces



Implementation (TUN → Socket)

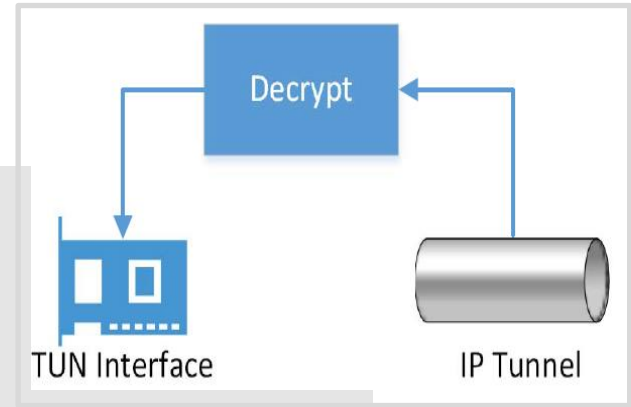
```
void tunSelected(int tunfd, int sockfd){  
    int len;  
    char buff[BUFF_SIZE];  
  
    printf("Got a packet from TUN\n");  
  
    bzero(buff, BUFF_SIZE);  
    len = read(tunfd, buff, BUFF_SIZE);  
    sendto(sockfd, buff, len, 0, (struct sockaddr *) &peerAddr,  
           sizeof(peerAddr));  
}
```



Note: the encryption step is omitted from the code (for the sake of simplicity)

Implementation (Socket → TUN)

```
void socketSelected (int tunfd, int sockfd){  
    int len;  
    char buff[BUFF_SIZE];  
  
    printf("Got a packet from the tunnel\n");  
  
    bzero(buff, BUFF_SIZE);  
    len = recvfrom(sockfd, buff, BUFF_SIZE, 0, NULL, NULL);  
    write(tunfd, buff, len);  
}
```



Note: the decryption step is omitted from the code (for the sake of simplicity)

Implementation (Monitoring the 2 Interfaces)

```
int main (int argc, char * argv[]) {
    int tunfd, sockfd;


    tunfd = createTunDevice();
    sockfd = connectToUDPServer();

    // Enter the main loop
    while (1) {
        fd_set readFDSet;

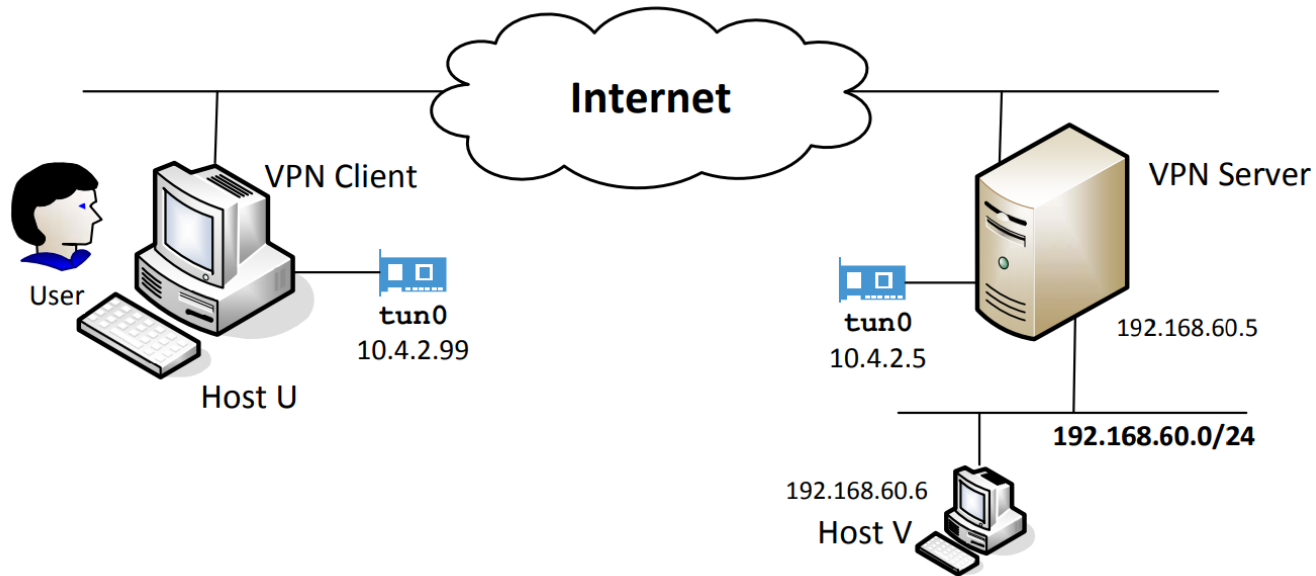
        FD_ZERO(&readFDSet);
        FD_SET(sockfd, &readFDSet);
        FD_SET(tunfd, &readFDSet);
        select(FD_SETSIZE, &readFDSet, NULL, NULL, NULL);

        if (FD_ISSET(tunfd, &readFDSet)) tunSelected(tunfd, sockfd);
        if (FD_ISSET(sockfd, &readFDSet)) socketSelected(tunfd, sockfd);
    }
}
```

select() will be blocked until one of the interfaces has data.



Case Study: Configuring a VPN



Configure VPN Server

- On VPN Server, we first run the server program.
- Configure the tun0 interface.
 - We use 10.4.2.0/24 as IP prefix for the TUN interface (for both VPN Client and VPN Server)
- The following two commands assign the IP address to the tun0, bring it up and then add a corresponding route to routing table.

```
$ sudo ifconfig tun0 10.4.2.5/24 up  
$ sudo route add -net 10.4.2.0/24 tun0
```


Configure VPN Client

- On VPN Client, we first run the client program.
- Add route for the 10.4.2.0/24 network.
- Add a route, so that all the packets for 192.168.60.0/24 are routed to the tun0 interface.

```
$ sudo ifconfig tun0 10.4.2.99/24 up  
$ sudo route add -net 10.4.2.0/24 tun0  
$ sudo route add -net 192.168.60.0/24 tun0
```

Configure Host V

- The reply packets should go back via the same VPN tunnel, so that they are protected.
- To ensure that, route all packets for the 10.4.2.0/24 network toward the tunnel.
- For Host V, we route such packets to VPN Server.
- Add the following routing entry to Host V:

```
$ sudo route add -net 10.4.2.0/24 gw 192.168.60.5 eth1
```

Testing VPN: ping Testing

- Ping Host V from Host U and we see the following result:

```
seed@User(10.0.2.6):$ ping 192.168.60.6
PING 192.168.60.6 (192.168.60.6) 56(84) bytes of data.
64 bytes from 192.168.60.6: icmp_req=1 ttl=63 time=2.41 ms
64 bytes from 192.168.60.6: icmp_req=2 ttl=63 time=1.48 ms
```

- The following figure shows the packets generated when we ping Host V (192.168.0.6).

No.	Source	Destination	Protocol	Length	Info
1	10.4.2.99	192.168.60.6	ICMP	100	Echo (ping) request id=0x0e85, seq=1/256, ttl=64
2	10.0.2.6	10.0.2.5	UDP	128	Source port: 59793 Destination port: 55555
3	10.0.2.5	10.0.2.6	UDP	128	Source port: 55555 Destination port: 59793
4	192.168.60.6	10.4.2.99	ICMP	100	Echo (ping) reply id=0x0e85, seq=1/256, ttl=63
5	10.4.2.99	192.168.60.6	ICMP	100	Echo (ping) request id=0x0e85, seq=2/512, ttl=64
6	10.0.2.6	10.0.2.5	UDP	128	Source port: 59793 Destination port: 55555
7	10.0.2.5	10.0.2.6	UDP	128	Source port: 55555 Destination port: 59793
8	192.168.60.6	10.4.2.99	ICMP	100	Echo (ping) reply id=0x0e85, seq=2/512, ttl=63

Testing VPN: telnet Testing

- The following result shows that we can successfully connect to the telnet server on Host V inside the private network

```
seed@User(10.0.2.6):$ telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 12.04.2 LTS
ubuntu login: seed
Password:
.....
seed@HostV(192.168.60.6):$
```

← **successfully logged in!**

Testing VPN: telnet Testing


- Let us break the tunnel and see what happens

Source	Destination	Protocol	Length	Info
10.0.2.17	10.0.2.18	UDP	243	Source port: personal-agent Destination port: 33025
10.0.2.18	10.0.2.17	UDP	96	Source port: 33025 Destination port: personal-agent
10.0.2.18	10.0.2.17	UDP	97	Source port: 33025 Destination port: personal-agent
10.0.2.17	10.0.2.18	ICMP	125	Destination unreachable (Port unreachable)
10.0.2.18	10.0.2.17	UDP	97	Source port: 33025 Destination port: personal-agent
10.0.2.17	10.0.2.18	ICMP	125	Destination unreachable (Port unreachable)
10.0.2.18	10.0.2.17	UDP	97	Source port: 33025 Destination port: personal-agent
10.0.2.17	10.0.2.18	ICMP	125	Destination unreachable (Port unreachable)
10.0.2.18	10.0.2.17	UDP	97	Source port: 33025 Destination port: personal-agent
10.0.2.17	10.0.2.18	ICMP	125	Destination unreachable (Port unreachable)
10.0.2.18	10.0.2.17	UDP	97	Source port: 33025 Destination port: personal-agent
10.0.2.17	10.0.2.18	ICMP	125	Destination unreachable (Port unreachable)

Observation: the telnet connection is not broken. TCP will keep resending packets, but they cannot be delivered because the tunnel is broken. Whatever we type in telnet will be buffered by TCP, not lost, but we can't see anything. As soon as we reconnect the tunnel, everything that we have typed will show up.

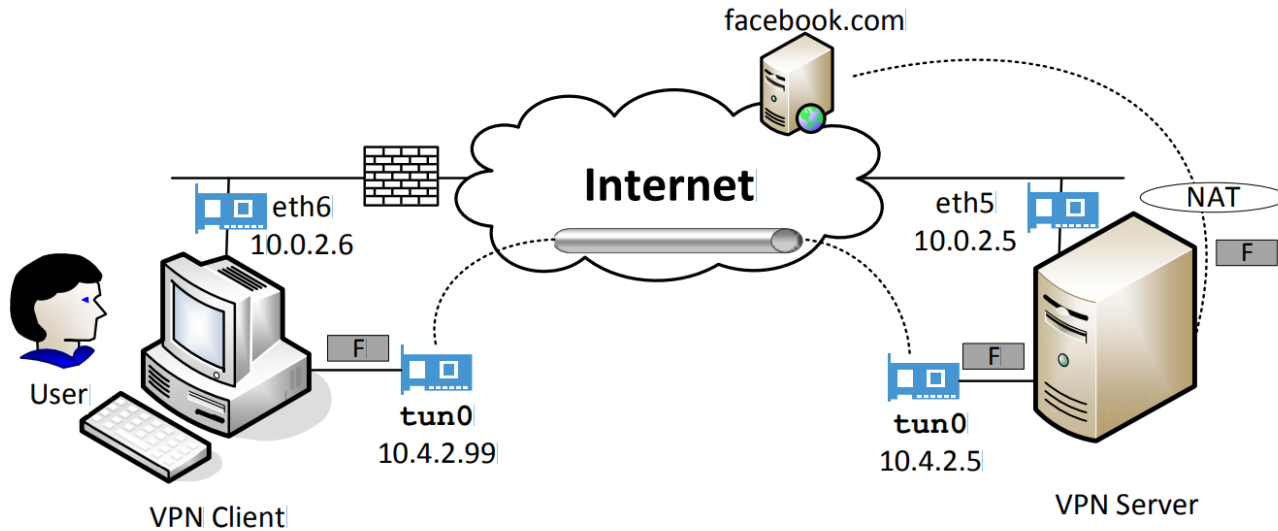
Bypassing Firewalls using VPN

Bypassing Firewall using VPN: the Main Idea



- Send our Facebook-bound packets to the TUN interface towards VPN server
- VPN server will release our Facebook-bound packets to the Internet
- Facebook's reply packets will be routed to the VPN server (**question: why**)
- VPN server sends the reply packets back to us via the tunnel

Experiment: Network Setup



Setting UP Firewall

- Setup firewall to block User from accessing Facebook
- We run the following command to get the list of IP prefixes owned by Facebook:

```
$ whois -h whois.radb.net -- '-i origin AS32934'
```

- We can also get IP addresses returned by Facebook's DNS server by running the following command (this IP address can change):
`dig www.facebook.com`

Blocking Facebook

One of the IP prefixes belong to Facebook

```
$ sudo ufw enable
$ sudo ufw deny out on eth6 to 31.13.0.0/16
$ sudo ufw status
Status: active
```

To	Action	From
--	-----	----
31.13.0.0/16	DENY OUT	Anywhere on eth6

Facebook becomes unreachable

```
seed@User(10.0.2.6):~$ ping www.facebook.com
PING star-mini.c10r.facebook.com (31.13.71.36) 56(84) bytes of data.
ping: sendmsg: Operation not permitted
ping: sendmsg: Operation not permitted
ping: sendmsg: Operation not permitted
ping: sendmsg: Operation not permitted
ping: sendmsg: Operation not permitted
```

Bypassing the Firewall

- We add a routing entry to the user machine, changing the route for all Facebook traffic. Instead of going through eth6, we use the TUN interface:

```
$ sudo route add -net 31.13.0.0/24 tun0
```

- The Facebook-bound packets are going through our tunnel.
- The Facebook-bound packets are hidden inside a packet going to the VPN server, so it does not get blocked.
- VPN server will release the packet to the Internet.
- Replies from Facebook will come back to VPN server, which will forward it back to us via the tunnel.

Summary

- What is VPN?
- IP tunneling
- IP tunneling using TLS/SSL
 - TUN/TAP interface
- Building a VPN using TUN/TAP interface
- Using VPN to bypass firewalls